

# Report For CS382 Homework 1

## Implementation of Smoothing Algorithm for N-gram Language Model

Yanming Liu; ID: 518030910393

October 17, 2021

**Declare** All the implementation here are finished by myself without the help of open source tools.

### 1 Smoothing Algorithm

#### 1.1 Hierarchical Dirichlet Language Model[1]

The hierarchical Dirichlet language model is a smoothed n-gram language model proposed by MacKay et al (1995)[1]. The article mainly discusses the prior and posterior distribution of n-gram language model parameters, and points out that the prior distribution can be a Dirichlet distribution.

Denote the language model parameters as  $\mathbf{p} = (p_1, p_2, \dots, p_n)$ , (here  $p_i = Pr(\text{some token} | \text{some context})$ ) before seeing the training data we have a prior distribution of  $\mathbf{p}$ :

$$P(\mathbf{p} | \alpha \mathbf{m}) = \frac{1}{Z(\alpha \mathbf{m})} \prod_{i=1}^n p_i^{\alpha m_i - 1} \delta(\sum_i p_i - 1) \equiv \text{Dirichlet}^{(n)}(\mathbf{p} | \alpha \mathbf{m}) \quad (1)$$

Here  $Z(\alpha \mathbf{m})$  is a normalizing constant (can be calculated with explicit formulas),  $\delta(\cdot)$  is Dirac delta function,  $\mathbf{m} = \mathbb{E}[\mathbf{p}]$  (so,  $\sum_i m_i = 1$ ) and  $\alpha$  is a positive scalar.

As far as I understand it, one of the advantages of using Dirichlet distribution as the prior distribution is that we can easily get a brief posterior distribution, after seeing the training corpus and getting the statistical information  $\mathbf{F} = (F_1, F_2, \dots, F_n)$ , where  $F_i$  means the number of tokens group  $i$ :

$$\begin{aligned} P(\mathbf{p} | \mathbf{F}, \alpha \mathbf{m}) &= \frac{P(\mathbf{p}, \mathbf{F}, \alpha \mathbf{m})}{P(\mathbf{F}, \alpha \mathbf{m})} = \frac{P(\mathbf{F} | \mathbf{p}) P(\mathbf{p} | \alpha \mathbf{m})}{P(\mathbf{F} | \alpha \mathbf{m})} \\ &= \frac{1}{P(\mathbf{F} | \alpha \mathbf{m}) Z(\alpha \mathbf{m})} \prod_i^n p_i^{F_i} \prod_i^n p_i^{\alpha m_i - 1} \delta(\sum_i p_i - 1) \\ &= \frac{1}{Z(\mathbf{F}, \alpha \mathbf{m})} \prod_i^n p_i^{F_i + \alpha m_i - 1} \delta(\sum_i p_i - 1) = \text{Dirichlet}^{(n)}(\mathbf{p} | \mathbf{F} + \alpha \mathbf{m}) \end{aligned} \quad (2)$$

Just change the Dirichlet distribution parameters from  $\alpha \mathbf{m}$  to  $\mathbf{F} + \alpha \mathbf{m}$ . And according to a good property of Dirichlet distribution,  $\mathbb{E}[p_i] \propto F_i + \alpha m_i$ , so we have:

$$\mathbb{E}[p_i] = \frac{F_i + \alpha m_i}{\sum_{i'} F_{i'} + \alpha m_{i'}} \quad (3)$$

So, when predict, we use the predictive distribution:

$$P(i | j, \mathbf{F}, \alpha \mathbf{m}) = \frac{F_{i|j} + \alpha m_i}{\sum_{i'} F_{i'|j} + \alpha m_{i'}} = \frac{F_{i|j} + \alpha m_i}{F_j + \alpha} = \frac{F_j}{F_j + \alpha} \frac{F_{i|j}}{F_j} + \frac{\alpha}{F_j + \alpha} m_i \quad (4)$$

If we let  $\lambda_j = \frac{\alpha}{F_j + \alpha}$ , and  $f_{i|j} = F_{i|j} / F_j$ , we get

$$P(i | j, \mathbf{F}, \alpha \mathbf{m}) = (1 - \lambda_j) f_{i|j} + \lambda_j m_i \quad (5)$$

Therefore, when we assume that the prior distribution of the parameters is Dirichlet distribution, we can derive the following:

- Posterior distribution is similar (almost the same) to the one of interpolation method.
- The interpolation coefficient  $\lambda_j = \alpha/(F_j + \alpha)$  can be calculated theoretically and not need to be tuned on dev set. And this result confirms a conclusion from experiments: the tokens group with the same frequency is better to use the same lambda.
- Explicit assumption of the prior distribution allows us to even calculate  $\alpha$  and  $\mathbf{m}$ .

It can be derived by Bayes formula that the optimal parameters  $\alpha$  and  $\mathbf{m}$  maximize  $P(F|\alpha\mathbf{m})$ . And we can calculate  $\alpha, \mathbf{m}$  by analysis method (the deriving process is complicated, we omit it here): Let  $N_{Fi}$  be the number of contexts  $j$  such that  $F_{i|j} \geq F$ , compute (when  $f$  is big enough,  $N_{fi}=0$ , so only need to calculate finite items):

$$G_i = \sum_{f=2}^{\infty} \frac{N_{fi}}{f-1}; \quad H_i = \sum_{f=2}^{\infty} \frac{N_{fi}}{(f-1)^2} \quad (6)$$

We define  $K(\alpha) = \sum_j \log \left( \frac{F_j + \alpha}{\alpha} \right) + \frac{F_j}{2\alpha(F_j + \alpha)}$

The optimal parameters  $\mathbf{u}^* = \alpha^* \mathbf{m}^*$  should satisfy:

$$u_i^* = \frac{2N_{1i}}{K(\alpha^*) - G_i + \sqrt{(K(\alpha^*) - G_i)^2 + 4H_iN_{1i}}} \quad (7)$$

And we know,  $\sum_i m_i = 1$ , so  $\alpha = \sum_i u_i$ . So far, we can calculate  $\alpha, \mathbf{u}$  in an iterative way: calculate  $\mathbf{u}$  from  $\alpha$  according to Eq.7 and then  $\alpha \leftarrow \sum_i u_i$ .

## 1.2 Tuned Interpolation Language Model

This model is an intuitive one. From **section 1.1** we can see that the smoothed distribution should be in such a family of distribution:

$$P(w_i | w_{i-n+1}^{i-1}) = (1 - \lambda_{w_{i-n+1}^{i-1}}) f_{i|j} + \lambda_{w_{i-n+1}^{i-1}} P(w_i | w_{i-n+2}^{i-1})$$

Here  $f_{i|j}$  is directly from the statistical frequency,  $P(w_i | w_{i-n+2}^{i-1})$  is used as the prior distribution, that is,  $m_{w_i}$  in Eq.5. We do not calculate  $\lambda$  explicitly here, instead we learn it from the dev set with maximum likelihood estimation:

$$\log P(F^{dev}; \lambda) = \sum_{j \in \{w_{i-n+1}^{i-1}\}} \sum_{i \in \{w_i\}} F_{i|j}^{dev} \log \left( (1 - \lambda_j) f_{i|j}^{train} + \lambda_j P(i|j^{backoff}) \right) \quad (8)$$

$$\frac{\partial}{\partial \lambda_j} \log P(F^{dev}; \lambda) = \sum_i F_{i|j}^{dev} \frac{P(i|j^{backoff}) - f_{i|j}^{train}}{(1 - \lambda_j) f_{i|j}^{train} + \lambda_j P(i|j^{backoff})} \quad (9)$$

It is obviously that  $\partial^2 \log P / \partial \lambda_j^2 < 0$ , so  $\partial \log P / \partial \lambda_j$  is monotonous about  $\lambda_j$ , so the optimal  $\lambda^*$  can be calculated by solving the equation  $\partial \log P / \partial \lambda_j = 0$  with bisection.

Inspired by the method in **section 1.1**, tokens group  $j$  with the same frequency  $F_j$  uses the same  $\lambda_{F_j}$  as parameter. It solves the following problems:

- Overfitting (good performance on the dev set but poor performance on the test set).
- When a tokens group  $j$  does not appear in the dev set, we can not assign a proper value for  $\lambda_j$  with Eq.9.

## 2 Experiment

### 2.1 Experiment Settings

- Model: I test the two models mentioned in **Section 1**.
- Model parameters: I test  $n = 2$  and  $n = 3$ .
- Corpus: Use the given corpus train\_set.txt, dev\_set.txt and test\_set.txt.
  - For model in **Section 1.1**, since it uses only train set and no need for dev set, so I use both train\_set.txt and dev\_set.txt as the train corpus, and use test\_set.txt as the test corpus.

- For model in **Section 1.2**, I use train\_set.txt as the train corpus, that is, do counting on train\_set.txt, use dev\_set.txt as the dev corpus, that is, do tuning on dev\_set.txt, and use test\_set.txt as the test corpus.
- Vocabulary: All the tokens appearing in the train corpus including ‘< s >’ and ‘< /s >’ and an additional ‘< unk >’ token.
  - That is to say, here I did not use low-frequency words as the unknown. Instead, I used words that never appeared in train corpus as the unknown. In order to let the model see ‘< unk >’ during counting, I manually insert  $n$  ‘< unk >’ into the train corpus, which will not affect the correctness of the model, but will only slightly hurt performance.
- Iteration accuracy:
  - For model in **Section 1.1**, in order to ensure that the conditional distribution predicted by model is normalized, we set  $\epsilon = 1e - 6$  when calculate  $\alpha$  and  $\mathbf{m}$ .
  - For model in **Section 1.2**, when I tune the parameter  $\lambda$ , I use  $\epsilon = 1e - 3$ .

## 2.2 Experiment Results

Table 1: Perplexity

	2-gram ppl	3-gram ppl
MacKay Method ( <b>Section 1.1</b> )	571	760
Simple Interpolation ( <b>Section 1.2</b> )	527	421

The perplexity of the two models under different parameters  $n$  is shown in Table 1. The result generally meets my expectation.

MacKay method, i.e. the model in **Section 1.1**, is beautiful in theory, but in fact its performance is not very good (maybe my implementation is poor). The original paper uses a lot of analysis methods and approximations in order to get a beautiful explicit results, so the actual result may not be as good as theoretically expected.

The simple interpolation method, i.e. the one in **Section 1.2**, is just to maximum the likelihood on dev corpus (equivalent to minimizing ppl). But I found in experiments that the optimization effect is not very significant. Compared with a random  $\lambda$ , the optimal  $\lambda$  can only reduce the ppl by less than 10%.

I have compared my result with that of my classmate who used a similar algorithm, and I found that my ppl is slightly larger than his. However, I double-checked my implementation and did not find any obvious errors. And I have checked that the conditional distribution reported by my model is normalized.

## References

- [1] David JC MacKay and Linda C Bauman Peto. A hierarchical dirichlet language model. *Natural language engineering*, 1(3):289–308, 1995.